

INDUSTRY PAPER

One Platform, Many Sources

The benefits and savings of unifying multi manufacturer and SCADA flow monitoring data

PREPARED BY

Environmental Data Services

DATE

October 2026

PLATFORM

FlowSense

Abstract

Australian water utilities typically hold sewer and stormwater flow data across three to six separate systems: one vendor portal per logger brand, the corporate SCADA historian, spreadsheets exported by consultants, and the files left behind by temporary monitoring programs. The result is a network that is monitored but not understood. Engineers spend their time moving data between systems rather than acting on it, and the questions that matter most, such as where inflow and infiltration enters the network or which pump station is running outside its duty, cannot be answered without stitching sources together by hand.

This paper sets out the operational and financial case for consolidating monitoring data from multiple instrument manufacturers and SCADA systems onto a single manufacturer agnostic platform. It identifies six cost categories that fragmentation creates, quantifies them for a representative regional utility with 40 permanent sites and a rolling temporary monitoring program, and describes the operational capabilities that only become possible once data sits in one place with one data model. The financial case rests on Australian evidence: recent NSW court penalties of \$200,000 to \$865,625 for single overflows, Sydney Water's reported \$400,000 a month in avoided incident costs from a two manufacturer sensor program, its finding that source control is 17 times cheaper than capacity works, and Unitywater's integration of four SCADA brands into one control room at no additional cost. Using FlowSense as a worked example, the paper shows that the largest savings do not come from licence consolidation but from overflows caught, capital avoided and engineering hours recovered.

1. Introduction: a network monitored in pieces

Most Australian sewer networks are monitored by instruments from several manufacturers, accumulated over 15 to 20 years of separate purchasing decisions. A permanent flow monitoring fleet bought under one contract, temporary loggers deployed by a consultant for an inflow and infiltration study, level sensors installed after an overflow incident, and pump station telemetry on the corporate SCADA system all measure the same network. Each arrives with its own portal, its own file format, its own alarm logic and its own login.

This is not the result of poor planning. It is the natural outcome of buying the right instrument for each job over a long period, in an industry where every logger manufacturer has historically bundled its hardware with a proprietary cloud portal. The instrument is fit for purpose. The data landscape it leaves behind is not.

The problem is well recognised. Industry commentary from the Australian Water Association notes that utility data is routinely scattered across SCADA platforms, asset systems, spreadsheets, partner reports and legacy databases, and that the less visible work of integration and governance is the real barrier to using it.¹⁰ Stantec observed in 2022 that getting SCADA to talk to other systems has historically been expensive, and that the sector has been moving toward feeding every system into one data store instead.¹¹

This paper takes that general observation and applies it to the specific case of sewer and stormwater flow monitoring. The question it answers is practical: what does a utility gain, in engineering hours, in dollars and in operational capability, when loggers from several manufacturers and the SCADA historian all report into one manufacturer agnostic platform? The audience is the water engineer or asset manager who owns the monitoring budget and has to justify a platform decision to a finance team.

Throughout, FlowSense is used as the worked example because it was built for exactly this problem: ingesting data from EDS's own Hawk and E-Flow instruments, from third party logger brands including Detectronic, MetaLog, Kallipr, Halytech, Ovarro, Metasphere and Unidata, from SCADA platforms including Ignition, Citect, Geo SCADA and AVEVA, and presenting all of it in one data model. The argument, however, applies to any platform designed on the same principle.

2. The real cost of fragmentation

Fragmentation rarely appears as a line in the budget. It shows up as engineering hours, slow responses and decisions made on partial data. The six categories below are where it hides. Each one is familiar to anyone who has tried to answer a simple question such as "what was the flow at every site in the catchment during last Tuesday's storm" when the sites report to three different portals.

Cost category	How it shows up in practice	Who carries it
Manual data handling	Exporting CSVs from each portal, aligning timestamps and units, merging in a spreadsheet before any analysis can start	Engineers and analysts
Duplicated licences and hosting	A subscription per manufacturer portal, often priced per logger, plus SCADA historian tags for the same sites	Operations budget
Training and access management	Every portal has its own login, its own alarm configuration and its own conventions; new staff must learn each, and departing staff leave accounts behind	Operations and IT
Slow fault detection	A stuck sensor or dropped channel is only noticed when someone opens that particular portal, which may be weeks after the fault	Data quality, later rework
Blind spots between systems	Pump station run data in SCADA and gravity flow data in a logger portal are never seen together, so the relationship between them is never analysed	Network planning
Vendor lock in	Once data lives in a manufacturer's portal, replacing that manufacturer means losing the history or paying to migrate it; the instrument decision becomes a platform decision	Procurement

The first category is the largest in dollar terms and the hardest to see. A utility engineer who spends two hours a week pulling data out of portals, cleaning it and lining it up is spending around 100 hours a year, before any analysis is done. At a fully loaded cost of roughly \$150 per hour (an approximate figure; substitute the utility's own rate), that is \$15,000 a year per engineer on data handling alone. For a team of four, with a consultant doing the same thing on the temporary monitoring program, the figure comfortably exceeds \$80,000 a year and produces nothing the utility can act on.

Stormwater ingress can lift sewer flows to 5 to 10 times dry weather levels.⁶ A velocity sensor that silently stops reading above a certain depth corrupts exactly the storm events a monitoring program exists to capture.

The fourth category is the most expensive in consequence. Sewer flow data is used to size augmentations, calibrate hydraulic models and locate inflow and infiltration. A velocity sensor that silently stops reading above a certain depth, or a logger that drops channels overnight, corrupts exactly the storm events the program exists to capture. If the fault is not seen until the next quarterly review, the storm season is lost and the model is calibrated on bad data.

The sixth category is strategic. The experience of the City of Fort Smith, which ran five separate SCADA systems installed over many years, is instructive: different software versions, different naming conventions, data trapped at each facility, and multiple integrators to manage.¹² Consolidating onto one platform cost the department two thirds of what upgrading the five legacy systems would have cost, and gave it a single view across four plants. Logger portals create the same trap on a smaller scale, one manufacturer at a time.

3. What a manufacturer agnostic platform looks like

The defining feature is not the dashboard. It is the data model underneath it: every site, channel and reading, regardless of which instrument produced it, is stored in one structure with one set of conventions for units, timestamps, quality flags and access. Once that is true, every capability built on top of it works for every site.




Figure 1. One data model turns four separate feeds into one view of the network.

Each feed on the left arrives in its own format: a logger's native protocol, a vendor API, a file drop, or an OPC or Modbus read from the SCADA historian. The platform's job is to map each one onto the same site and channel structure so that a depth reading from a third party logger and a depth reading from an EDS Hawk are indistinguishable downstream.

Four properties separate a true multi manufacturer platform from a dashboard that merely links to other portals:

1. **Native ingestion, not screen scraping.** Data arrives through the manufacturer's API or protocol at the logger's native interval, with the original quality flags preserved. FlowSense, for example, ingests

Detectronic loggers through DetectData and is adding Kallipr sensors alongside EDS's own Hawk loggers.

2. **One site, one history.** When an instrument at a site is replaced with another brand, the site's history continues unbroken. The instrument is an attribute of the channel, not the owner of the data.
3. **SCADA as a peer, not a destination.** Pump station run times, wet well levels and flow meter totals from the historian sit beside gravity main flow data in the same time series store, so a pump station's behaviour can be compared with the flows arriving at it.
4. **The utility owns the data.** Export to CSV, to a hydraulic model format or to the utility's own data warehouse is a standard function, not a professional services engagement.

The last property matters most for procurement. If the platform holds the history and exports it freely, the utility can buy the best instrument for each site, from any manufacturer, without changing anything else.

Manufacturers and systems FlowSense ingests

FlowSense is deliberately open about which brands it reads. It takes data in by six routes, listed below with the companies and products connected through each as at October 2026. Many appear under more than one route because the choice depends on what the client's own system can do and what their IT policy allows.

EDS NATIVE EDS instruments report directly

Configured in EMS-Flow at commissioning; no integration step.

EQUIPMENT EDS Hawk 4G logger, E-Flow 3, AVFS, MetaFlow MF-290

API POLLING FlowSense polls the vendor's cloud

The platform calls the client's HTTP API on a schedule and reads what has arrived since last time. Nothing is installed and nothing is opened on their side.

CLOUDS Detectronic DetectData, MetaLog m2m, Kallipr Kloud, Xylem Vue Secure Connect, eagle.io, Taggle Aqualus, Metasphere Palette, Unidata Neon, AVEVA Insight, VTScada REST, Kisters Hydstra and WISKI, Aquarius, NNNCo N2N-DL, Thinxtra, Myriota, Ellenex

EQUIPMENT Detectronic MSFM4 and LIDoTT, MetaLog loggers, Kallipr Captis, Taggle meters, Metasphere Point Orange, Unidata Neon, Ellenex, Myriota, LoRaWAN sensors on NNNCo or Thinxtra

HTTP PUSH Their SCADA pushes to FlowSense

The client's system posts readings to a secure endpoint with a write only token. Nothing is opened on their network.

SCADA Ignition, VTScada, Geo SCADA (ClearSCADA), AVEVA System Platform, Plant SCADA (Citect) via Cicode, Node-RED

EQUIPMENT Halytech Spider, Ovarro TBox, The Things Stack, ChirpStack, LoRaWAN servers with webhooks

INTEGRATORS SAGE Automation, SAFEgroup Automation, ACE, Beetle Engineering

MQTT FlowSense subscribes to their broker

The platform connects out to the client's MQTT broker and listens on their topics. Sparkplug B is read on its own namespace, so Ignition MQTT Transmission needs nothing added.

- SCADA** Ignition with MQTT Transmission (Sparkplug B or JSON), Node-RED, Mosquitto, EMQX, HiveMQ
- EQUIPMENT** Kallipr Captis, Ovarro TBox, Schneider SCADAPack 470 and 570, NNNCo N2N-DL
- INTEGRATORS** SAGE Automation, SAFEgroup Automation

FILE DROP They push files to an EDS hosted SFTP folder

The client's system uploads CSV exports on a schedule they own. Nothing is exposed at their end. This is the usual answer for councils that can export a file but will not publish a server for an outside platform to collect from.

- SCADA** Geo SCADA scheduled upload, Plant SCADA (Citect) scheduled report upload, AVEVA Historian scheduled export, Kisters Hydstra export
- SCRIPTS** WinSCP scheduled task, PowerShell Posh-SSH, sftp batch file, Python paramiko
- INTEGRATORS** SAFEgroup Automation, SAGE Automation, Aquamonix Solutions

FILE PICKUP FlowSense collects their file exports

The platform polls a folder on the client's SFTP or FTPS server and ingests new CSV exports.

- SCADA** Plant SCADA (Citect) reports, Geo SCADA CSV export, AVEVA Historian (Wonderware), WinCC OA, Hydstra
- EQUIPMENT** Halytech Spider, Unidata Neon, Metasphere Point Orange, Ovarro TBox, Campbell Scientific LoggerNet
- INTEGRATORS** Aquamonix Solutions, ACE, Radtel

The SCADA routes are two way: FlowSense also forwards validated data back to SCADA, telemetry and reporting systems in the formats they already accept, so it drops into an existing stack without the client changing anything downstream.

4. Operational benefits

The benefits below are listed in the order a utility usually notices them. The first two arrive within weeks of consolidation. The last three are only possible once data from different sources has been sitting side by side long enough to be compared.

One view of the network

An operator on a wet night sees every monitored site on one map, with every alarm in one list, regardless of which company made the instrument. The time saved on any single event is modest. The value is that nobody has to remember which portal a site lives in, and that the temporary loggers deployed last month are visible to the same people who watch the permanent fleet.

Consistent quality assurance

When every channel passes through the same validation rules, fault detection becomes systematic rather than incidental. Flat lining, stuck values, overnight dropouts, reporting gaps and out of range readings are flagged the same way for every brand. Experience auditing third party logger data for Australian clients shows that these faults are common and often go unnoticed for weeks in the manufacturer's own portal, because that portal is opened only when someone needs a specific export. A platform that checks every channel every day catches the fault on day one, while the site can still be visited before the next storm.

Alarms that reflect the network, not the instrument

A manufacturer portal can alarm on its own instrument's thresholds. A unified platform can alarm on relationships: a pump station whose run hours rise while the upstream gravity flow does not, a level site surcharging while the SCADA shows the downstream pump healthy, or a flow site reporting dry weather flow during a rain event the gauge recorded. These are the alarms that find blockages, failed pumps and illegal connections, and none of them can be configured when the data is split.

Inflow and infiltration analysis across the whole catchment

Inflow and infiltration studies depend on comparing wet weather response across sites in the same catchment during the same events. With data from a mix of permanent, temporary and SCADA sources in one time series store, the engineer can rank sub catchments by their wet weather response factor from the platform itself, without first building the merged dataset by hand. The same data then goes to the hydraulic modeller in one export, with the quality flags intact, which shortens model calibration and removes the argument about whose data is right.

Pump station performance from data the utility already has

SCADA historians hold years of pump run and wet well level data that is rarely analysed because it sits in a system engineers do not routinely use. Bringing it onto the monitoring platform allows pump station inflow to be derived from wet well drawdown and compared with the measured flows upstream, turning an existing data source into a flow monitoring site at no instrument cost. FlowSense includes a pump station flow feature built on exactly this method. For a utility with dozens of pump stations, this is the single largest expansion of network visibility available without buying hardware.

5. The financial case

The money is in four places: the cost of a spill that was not caught, the capital that good monitoring data lets a utility avoid, the integration and licence spend that consolidation removes, and the engineering time that fragmentation consumes. Each is supported below by Australian evidence first, with overseas cases where they add something. The worked example at the end of the section is labelled as an example; the figures before it are not.

What a missed overflow costs in Australia

An overflow that a connected alarm would have caught is the single most expensive consequence of fragmented monitoring. The NSW Land and Environment Court outcomes below are all recent, all for Sydney Water, and all public.

Incident	What happened	Outcome
Naremburn, October 2020	Partial collapse of a sewer built between 1916 and 1930; about 16 million litres overflowed into Flat Rock Creek over five days; 39 residents relocated	\$365,625 fine plus \$500,000 to environmental projects, total \$865,625 (NSW EPA, 2023) ¹
Strathfield, January 2022	Split in a rising main; 700,000 to 1.2 million litres of untreated wastewater reached Powells Creek, a tributary of the Parramatta River	\$200,000 to the NSW Environmental Trust after a guilty plea to failing to maintain equipment (NSW EPA, 2024) ²
Carramar, 2026	Sewage spill; guilty plea to two offences in the Land and Environment Court	\$240,000 penalty (NSW EPA, 2026) ³

The penalties are the visible part. In the Naremburn case the sewage ran for five days and the judge found Sydney Water could have investigated and repaired the asset earlier; the maximum penalty available was \$3 million.¹ Beyond the court, each event carries clean up, relocation, reporting and reputational costs that are not itemised publicly but that the utility pays in full.

The scale of the exposure is large. Sydney Water recorded 11,220 dry weather overflows from sewer main chokes in 2023 to 24, of which 487 reached a waterway.⁴ Every one of those is a potential prosecution if it reaches a sensitive receiver, and every one started as a rising level that a sensor could have reported.

What early detection is worth: Sydney Water's own numbers

Sydney Water's \$15 million wastewater IoT program is the largest in Australia and the clearest Australian evidence that detection pays. By October 2023 it had installed 15,000 level sensors, on its way to more than 26,000 across 4,600 km of network, in the areas that account for 70 per cent of overflows to waterways. The devices come from two manufacturers, Kallipr and Metasphere. The program detects about 20 blockages a month in the gravity network before they become overflows, and Sydney Water's IoT manager put the result at "\$400,000 a month in avoidance costs".⁵ At that rate the program's capital cost is recovered in a little over three years, before counting any avoided prosecution.

Sydney Water bought sensors from two manufacturers for one program. Detection only pays if every alarm, from every brand, lands in front of the same operator at the same time. That is the platform's job.

Two points follow for every smaller utility. First, the saving is proportional to the network, not to the size of the utility; a regional council with a tenth of the overflows and a tenth of the sensors should expect a tenth of the avoided cost, which is still tens of thousands of dollars a month. Second, a manufacturer agnostic platform is not an optional extra on a sensor program but the thing that turns the sensors into avoided cost.

Capital avoided: source control is 17 times cheaper than capacity

The largest financial number in Australian sewer monitoring comes from Sydney Water's Wet Weather Overflow Monitoring Program synthesis report. Meeting its overflow frequency targets through storage and transport capacity would cost more than \$18 billion in 2022 dollars. Source control, which means finding and

removing inflow and infiltration, was assessed as about 17 times cheaper, and is the approach the utility has adopted, with about \$1.5 billion invested since 1998.⁶ The 2025 IPART price determination funded a further \$243 million for wet weather overflow abatement over five years.⁷

Source control is only cheaper if the utility knows where the water is getting in, and that knowledge comes from flow monitoring compared across sites during the same storms. Coffs Harbour City Council's inflow and infiltration program is the regional version of the same finding: sewer load rising to more than 10 times average dry weather flow in rain, direct inflow from illegal connections consuming up to half the network's capacity, and an expectation of deferring and downsizing treatment plant upgrades once the inflow is removed.⁸ The data behind that decision came from pump station records and flow monitoring that had to be analysed together. A platform that already holds both is the difference between a six week study and a six month one.

What consolidation itself saves: integration, licences and capital

The Australian case closest to the subject of this paper is Unitywater in South East Queensland. Its 17 sewage treatment plants ran on four different SCADA brands and sat outside the 24 hour control room's view; after hours an on call operator relied on paged alarms and remote dial in. The Operational Technology team brought all 17 plants into the control room in six weeks using existing tools, at no additional cost, and in doing so deferred capital spending on bandwidth upgrades and additional SCADA licensing.⁹ The lesson is not the specific tools but the pattern: the integration cost that had been "prohibitive" was mostly the cost of doing it plant by plant and brand by brand.

Two overseas cases put numbers on the same pattern. The City of Fort Smith, Arkansas, replaced five separate SCADA systems with one platform for two thirds of the cost of upgrading the five, and gained a single terminal view across four plants.¹² Mountain Regional Water District in Utah used the network wide visibility from its consolidated platform to save about US\$400,000 a year on pumping energy, which its water and energy manager said funded his department and its SCADA upgrades.¹³

Engineering time: the cost nobody budgets

There is no Australian survey of how long water engineers spend moving monitoring data between systems. The closest evidence is from data professionals generally: Anaconda's 2020 survey of nearly 2,400 practitioners found data loading took 19 per cent of their time and data cleansing more than a quarter, so close to half of every working day went on getting data into a usable state before any analysis.¹⁴ Utility engineers working across three vendor portals and a SCADA historian are doing the same job with worse tools. The worked example below assumes two hours a week per engineer, well under that survey's figure.

Worked example for a regional utility

The figures in this example are assumptions, chosen to be conservative, for a utility with 40 permanent flow and level sites across two or three logger brands, a rolling temporary monitoring program run by a consultant, and around 20 pump stations on SCADA. They are here to show the shape of the savings. Replace every number with the utility's own.

- Four engineers and analysts each spend two hours a week (100 hours a year) exporting, cleaning and aligning monitoring data across portals, at a fully loaded cost of \$150 an hour.
- The consultant running the temporary program bills 160 hours a year on the same data handling, at a similar rate.

- Two significant storm events a year are lost to sensor faults that went unnoticed in a vendor portal, each costing \$20,000 in redeployment, extended monitoring and modelling rework.
- Two vendor portals at \$6,000 a year each are retired because the platform ingests those loggers directly.
- Six pump stations become usable flow monitoring sites from existing SCADA data instead of being instrumented at \$5,000 each, counted as a one off avoided cost spread over a year.
- Training and account administration drop by one week of staff time a year.

Engineering hours are the largest saving in a representative regional utility

Estimated annual saving, AUD, total \$174k a year




Figure 2. Illustrative figures for a utility with 40 permanent sites and a rolling temporary program. Substitute the utility's own rates and site counts.

On these assumptions the saving is about \$174,000 a year, and licence savings, which are what a finance team usually asks about first, are the smallest category. The example deliberately excludes the avoided overflow costs and the capital effects described above, because they depend on the individual network. For most utilities a single avoided prosecution, or a single deferred augmentation, exceeds the example's total several times over.

Platform cost

Against these savings sits the platform subscription and the one off integration effort for each new source. The subscription is a fraction of the engineering time figure alone, and integration is incurred once per manufacturer, not once per site, so the cost of adding a fourth brand to a platform that already reads three is small and known in advance.

6. Case examples

A mixed fleet on one platform: Toowoomba Regional Council

EDS has supplied, installed and maintained sewer flow monitoring for Toowoomba Regional Council continuously since 2005, and today maintains more than 23 permanent and temporary sites. Over two decades the fleet has naturally come to include instruments of different generations and makes. Bringing that fleet into one view, with the same validation applied to every site regardless of age or brand, is the practical test of the argument in this paper.

A council replacing a vendor dashboard with an open platform

A regional New South Wales council running Kallipr Captis loggers approached EDS in 2026 because the vendor dashboard was cumbersome for day to day use and gave no view of sensor health across the fleet. The integration path chosen was to poll the Kallipr cloud API into FlowSense, so the council keeps its existing hardware and gains fleet health classification, consistent validation and the same interface as its other monitoring.

Two manufacturers, one program: Sydney Water

Sydney Water's wastewater IoT program is the largest Australian example of the problem this paper addresses. More than 26,000 level sensors from two manufacturers, Kallipr and Metasphere, feed alarms to the utility's operations teams, and the value, about 20 blockages a month caught before they overflow, depends entirely on those alarms arriving in one place regardless of which device sent them.⁵ The program was built from the start on the principle that the sensor brand should not determine where the data goes.

Consolidation at a Queensland utility: Unitywater

Unitywater's 17 treatment plants ran on four SCADA brands and were invisible to its 24 hour control room until the Operational Technology team integrated all of them in six weeks, deferring bandwidth and SCADA licence spending in the process.⁹ It is the same pattern as this paper's argument, applied to treatment rather than network monitoring: the barrier was never the technology but the cost of doing the integration one brand at a time.

Published experience from overseas utilities

The pattern repeats wherever a utility has consolidated. The City of Fort Smith in Arkansas replaced five SCADA systems, installed at different times with different naming conventions and data trapped at each plant, with one platform, for two thirds of the cost of upgrading the five.¹² Mountain Regional Water District in Utah used the resulting visibility to save around US\$400,000 a year on energy.¹³ Neither utility bought new sensors to achieve this. The gain came from putting data they already collected in one place.

7. Implementation considerations

Consolidation is a procurement and governance decision as much as a technical one. Five questions decide whether it succeeds.

Who owns the data? The contract should state that monitoring data belongs to the utility, that full history can be exported in a documented format at any time, and that the platform provider has no claim on it if the relationship ends. This is the clause that turns the platform from a new silo into the end of silos.

Which protocols are supported, and who pays for new ones? A platform that reads two brands today should be able to read a third without a bespoke project. Ask for the list of supported manufacturers and transports (vendor APIs, MQTT, HTTP push, SFTP file drop, OPC or Modbus for SCADA) and for the commercial terms for adding a new one. Integration is incurred once per manufacturer; it should be priced that way.

How does SCADA stay secure? Reading the historian is a one way export from the operational network and should be designed as one, through a data diode, a DMZ historian replica or a scheduled export, never a direct connection from the cloud into the control network. Writing validated data back into SCADA is useful but is a separate decision for the OT team.

What happens to the existing portals? Most utilities run the old portals in parallel for one storm season before switching them off. This gives the operations team confidence in the new alarms and gives the platform provider a reference period to prove the validation rules catch what the old portals caught.

Who changes their routine? The savings described above are only realised if engineers stop opening the vendor portals. That takes a visible decision from the monitoring program owner, a short training session, and a review after the first wet season to confirm the manual exports have actually stopped.

8. Conclusion and recommendations

A utility that monitors its network through several manufacturer portals and a separate SCADA historian is paying for the same data several times over: once in licences, again in the hours spent moving it, and most expensively in the faults, storms and relationships it never sees. Consolidating onto a manufacturer agnostic platform removes the first two costs directly and makes the third visible for the first time.

The recommendations for a utility considering this step are short:

1. **Price one overflow.** Take the utility's last significant spill and add the clean up, reporting, regulator response and any penalty. Recent NSW outcomes ran from \$200,000 to \$865,625 per event. That figure is the benchmark every platform decision should be measured against.
2. **Count the hours.** Ask each engineer and the consultant how long they spend each week getting monitoring data into a usable form. The answer is usually the business case on its own.
3. **Separate the instrument decision from the platform decision.** Specify the platform on its ability to ingest any manufacturer and export everything, then buy instruments on merit for each site.
4. **Bring SCADA in as a data source early.** Pump station data the utility already holds is the cheapest expansion of network visibility available.
5. **Run the old portals in parallel for one wet season,** then switch them off and confirm the manual exports have stopped.

FlowSense was built on these principles. EDS welcomes enquiries from utilities and consultants who want to test the argument against their own fleet.

www.e-d-s.com.au

References

1. NSW Environment Protection Authority (2023). *Sydney Water to pay over \$865k after huge sewage overflow*. Media release, 11 July 2023. epa.nsw.gov.au
2. NSW Environment Protection Authority (2024). *Sydney Water to pay \$200,000 for untreated sewage spill*. Media release, December 2024. epa.nsw.gov.au
3. NSW Environment Protection Authority (2026). *Sydney Water cops \$240,000 penalty over Carramar sewage spill*. Media release, May 2026. epa.nsw.gov.au
4. Yan, Y. (2024). *Containment and removal of sewage overflows from impacted surface waters*. Water e-Journal, Vol 10 No 5, Australian Water Association.
5. Infrastructure Magazine (2023). *Early detection: how Sydney Water's IoT rollout is preventing overflows*. 14 November 2023. infrastructuremagazine.com.au
6. Sydney Water (2024). *Wet Weather Overflow Monitoring Program 2016 to 2024 Synthesis Report*. sydneywater.com.au
7. Sydney Coastal Councils Group (2025). Submission to the IPART review of Sydney Water prices from 1 July 2025. ipart.nsw.gov.au
8. Australian Water Association (2024). *Council manages inflow and infiltration to optimise operations and capital upgrades*. Water Source, sponsored content. awa.asn.au
9. Utility Magazine. *SCADA integration project delivers improved asset visibility*, Unitywater. utilitymagazine.com.au
10. Westphal, T. (2026). Australian Water Association commentary, summarised in *Water Utilities See Clear AI Use Cases, but Data Foundations Remain the Barrier*, Net Zero Compare.
11. Stantec (2022). *Water analytics data helping the rise of smart utilities*. stantec.com
12. Wastewater Digest (2021). *Replacing Five SCADA Systems with One*, City of Fort Smith, Arkansas (sponsored content). wwdmag.com
13. Industrial Ethernet Book (2018). *Increased flexibility and lower costs for water utilities*, Mountain Regional Water District and Park City Water, Utah. iebmedia.com
14. HPCwire (2020). *Data prep still dominates data scientists' time, survey finds*. Anaconda State of Data Science 2020. hpcwire.com
15. Environmental Data Services (2026). FlowSense platform documentation and project records (internal).

All dollar figures in the worked example are illustrative assumptions for a representative regional utility and should be replaced with the reader's own rates, site counts and contract values.